

Order-to-Cash (O2C) Business Analytics Platform

Cloud Data Engineering + Data Warehouse + BI Reporting

Executive Summary

Many companies generate revenue but still struggle with cash flow.

The real problem is not sales. The real problem is **how efficiently orders convert into collected payments.**

This project builds an end-to-end Order-to-Cash analytics platform that tracks the full financial lifecycle:

Customer → Order → Shipment → Invoice → Payment

The system ingests raw operational data, validates it, models it into an analytics warehouse, and exposes business insights through an interactive dashboard.

The goal is to help management answer:

- Are we growing?
 - Are customers actually paying?
 - Where is revenue getting stuck?
 - Which customers are valuable vs risky?
-

Business Problem

In many organizations, sales, logistics, billing, and finance teams operate in separate systems.

Because of this:

- Finance cannot track unpaid invoices quickly
- Delayed deliveries go unnoticed
- High-value customers are not identified
- Management sees revenue but not collection efficiency

The business requires a single analytics system that monitors the **entire revenue cycle**, not just sales.

This project simulates a company's operational data and builds a reporting platform to monitor financial performance and cash collection behavior.

2. Solution Overview

The project implements a cloud-based data pipeline and reporting system.

Raw CSV Files

- Azure Blob Storage
- Python ETL (Cleaning + Validation)
- Azure SQL Database (Analytics Warehouse)
- Power BI Dashboard

In a production environment, the pipeline would be scheduled and automated using **Azure Data Factory**. Azure Data Factory would orchestrate ingestion, trigger validation, and refresh the warehouse automatically. The current project runs manually to demonstrate logic and data quality controls, but the architecture is ready for automation.

Resources				
Recently viewed	Favorites	Type	Alerts	Last viewed
 otc-sql-server-01		SQL server	0	February 6, 2026
 otc_analytics.db		SQL database	0	January 31, 2026
 omstd		Storage account	0	January 17, 2026
 omazuredatastore01		Resource group	0	January 17, 2026

Technology Stack

Layer	Technology Used
Data Source	Raw CSV Operational Files
Cloud Storage	Azure Blob Storage
Data Processing	Python (Pandas ETL + Validation)
Data Warehouse	Azure SQL Database
Data Modeling	Star Schema (Facts & Dimensions)
Visualization	Power BI Dashboard

4. Data Ingestion (Cloud Storage)

All operational CSV files are stored in Azure Blob Storage as the raw data layer.

Tables ingested:

- Customers
- Orders
- Shipments
- Invoices
- Payments

Blob Storage acts as a landing zone similar to how companies receive ERP exports.

The screenshot shows the Azure Blob Storage interface. The left sidebar lists 'File shares', 'Queues' (which is selected), and 'Tables'. The main area shows a list of blobs under the 'raw-data' container. The list includes:

	Name	Last modified	Access tier	Blob type	Size	Lease state
0	customers_raw	17/01/2026, 13:20:58	Hot (Inferred)	Block blob	22.62 KiB	Available
0	invoices_raw	17/01/2026, 13:20:58	Hot (Inferred)	Block blob	157.2 KiB	Available
0	orders_raw	17/01/2026, 13:20:58	Hot (Inferred)	Block blob	307.01 KiB	Available
0	payments_raw	17/01/2026, 13:20:58	Hot (Inferred)	Block blob	158.67 KiB	Available
0	shipments_raw	17/01/2026, 13:20:58	Hot (Inferred)	Block blob	230.36 KiB	Available

5. ETL Pipeline & Data Validation

A Python ETL script was developed to extract data from Blob Storage and load it into Azure SQL.

The ETL process performs:

Data Cleaning

- Date conversion
- Null handling
- Data type correction
- Status normalization

Data Validation Rules

Customers

- No duplicate customer IDs
- Valid regions only

Orders

- Valid customer reference (FK validation)
- Valid order status
- Positive order values

Shipments

- Linked to valid order
- Delivery date cannot be before ship date
- Delivered shipments must have delivery date

Invoices

- Linked to valid order
- Invoice amount > 0
- Valid invoice status

Payments

- Linked to valid invoice
- Payment amount > 0
- Valid payment status

The ETL is **idempotent**.

Every run clears and reloads tables safely in dependency order.

Load Order:

Customers → Orders → Shipments → Invoices → Payments

In a production environment, Azure Data Factory would schedule and trigger this ETL automatically.

etl_pipeline_execution.png

```
C:\Users\Om\azure_otc_projects>python -u etl\otc_etl_with_validation.py
Clearing tables in FK-safe order
All tables cleared
Reading customers_raw.csv from Blob Storage
Validating customers data
Clearing existing customers table
Loading customers into Azure SQL
Customers load completed successfully
Reading orders_raw.csv from Blob Storage
Unique order_status values:
['Completed' 'Cancelled']
Rows with invalid promised_ship_date:
    order_id  customer_id      order_date  order_value order_status promised_ship_date
2            3        170  2025-09-03 06:24:32.323090    40389.30   Cancelled          NaN
11           12        46  2024-09-08 06:24:32.323113   41491.78   Cancelled          NaN
25           26        107 2024-08-02 06:24:32.323143    39827.04   Cancelled          NaN
31           32        292 2024-12-20 06:24:32.323156   15149.17   Cancelled          NaN
32           33        257 2025-05-12 06:24:32.323159    23229.87   Cancelled          NaN
...
3949          3950        137 2024-08-18 06:24:32.369225    24195.74   Cancelled          NaN
3954          3955        258 2025-06-05 06:24:32.369236    26593.79   Cancelled          NaN
3955          3956        58  2024-10-31 06:24:32.369238    38316.18   Cancelled          NaN
3980          3981        363 2025-07-06 06:24:32.369290    14742.06   Cancelled          NaN
3983          3984        324 2024-12-26 06:24:32.369296    43294.80   Cancelled          NaN
[396 rows x 6 columns]
Reading customers table for FK validation
Validating orders data
Loading orders into Azure SQL
Orders load completed successfully
Reading shipments_raw.csv from Blob Storage
Unique shipment_status values:
['Delivered' 'Delayed']
Reading orders table for FK validation
Validating shipments data
Loading shipments into Azure SQL
Shipments load completed successfully
Reading invoices_raw.csv from Blob Storage
Unique invoice_status values:
['Issued' 'Pending']
Reading orders table for FK validation
Validating invoices data
Loading invoices into Azure SQL
Invoices load completed successfully
Reading payments_raw.csv from Blob Storage
Unique payment_status values:
['Paid' 'Partial']
Reading invoices table for FK validation
Validating payments data
Loading payments into Azure SQL
Payments load completed successfully
C:\Users\Om\azure_otc_projects>
```

6. Data Warehouse Design (Azure SQL Analytics Layer)

After the ETL process, the cleaned data is stored in an **Azure SQL analytics warehouse**.

Power BI does NOT connect to raw CSV files or staging tables. It connects only to this analytics layer.

A **Star Schema** model was implemented to support reliable reporting and fast aggregations.

Why this was necessary:

Operational data contains multiple transactions (orders, shipments, invoices, payments).

If reporting is built directly on transactional tables:

- filters propagate incorrectly
- DAX measures calculate wrong totals
- performance becomes slow

The star schema fixes this by separating descriptive data from transactional events.

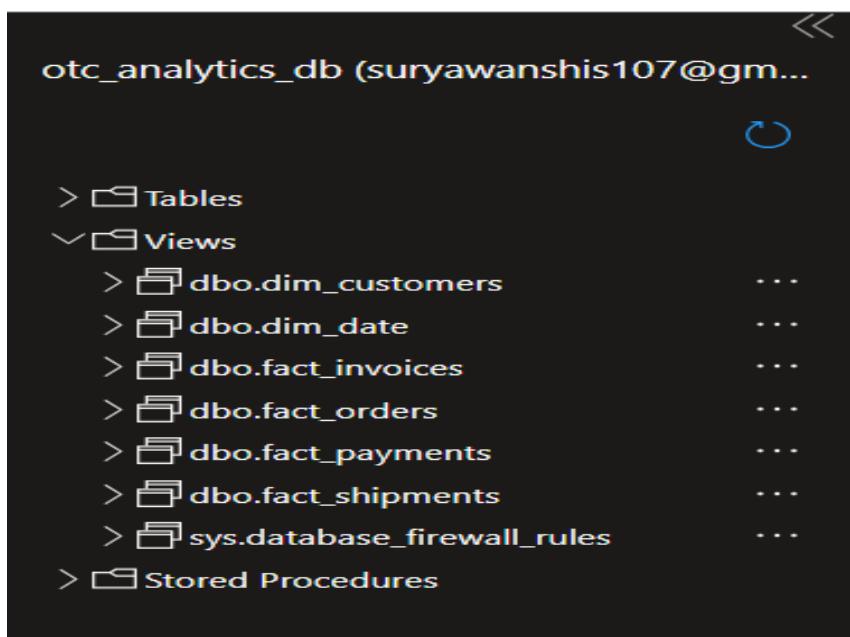
Schema Structure

Dimension Tables (used for filtering)

- **dim_customers** → customer and region information
- **dim_date** → calendar hierarchy (year, month, weekday)

Fact Tables (business events)

- **fact_orders** → order placed and order value
- **fact_shipments** → delivery execution
- **fact_invoices** → billing generation
- **fact_payments** → cash received



Business Process Relationship

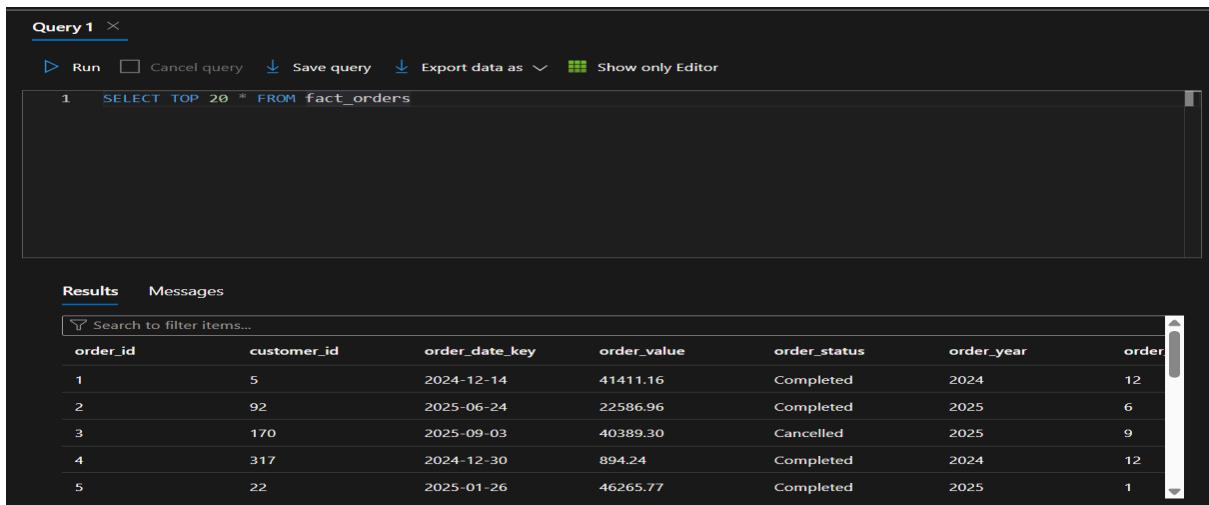
The warehouse models the real Order-to-Cash cycle:

Customer → Order → Shipment → Invoice → Payment

Relationships:

- fact_orders.customer_id → dim_customers.customer_id
- fact_orders.order_date_key → dim_date.date_key
- fact_shipments.order_id → fact_orders.order_id
- fact_invoices.order_id → fact_orders.order_id
- fact_payments.invoice_id → fact_invoices.invoice_id

This allows Power BI to trace revenue from customer acquisition to cash collection.



The screenshot shows the Power BI Query Editor interface. At the top, there are buttons for 'Run', 'Cancel query', 'Save query', 'Export data as', and 'Show only Editor'. Below the toolbar, a code editor window contains the following SQL query:

```
1  SELECT TOP 20 * FROM fact_orders
```

Below the code editor is a results grid. The columns are labeled: order_id, customer_id, order_date_key, order_value, order_status, order_year, and order_month. The data is as follows:

order_id	customer_id	order_date_key	order_value	order_status	order_year	order_month
1	5	2024-12-14	41411.16	Completed	2024	12
2	92	2025-06-24	22586.96	Completed	2025	6
3	170	2025-09-03	40389.30	Cancelled	2025	9
4	317	2024-12-30	894.24	Completed	2024	12
5	22	2025-01-26	46265.77	Completed	2025	1

Table Grain (Important for BI Calculations)

Each row represents a real business event:

Table	Represents
fact_orders	One order
fact_shipments	One shipment
fact_invoices	One invoice
fact_payments	One payment

Correct grain is required for accurate metrics like AOV, Collection Efficiency and CLV.

Why This Matters for Power BI

Power BI uses:

Dimensions → filtering

Facts → aggregation

Because filters flow one direction (dimension → fact):

- measures calculate correctly
- no ambiguous relationships
- dashboard performance improves

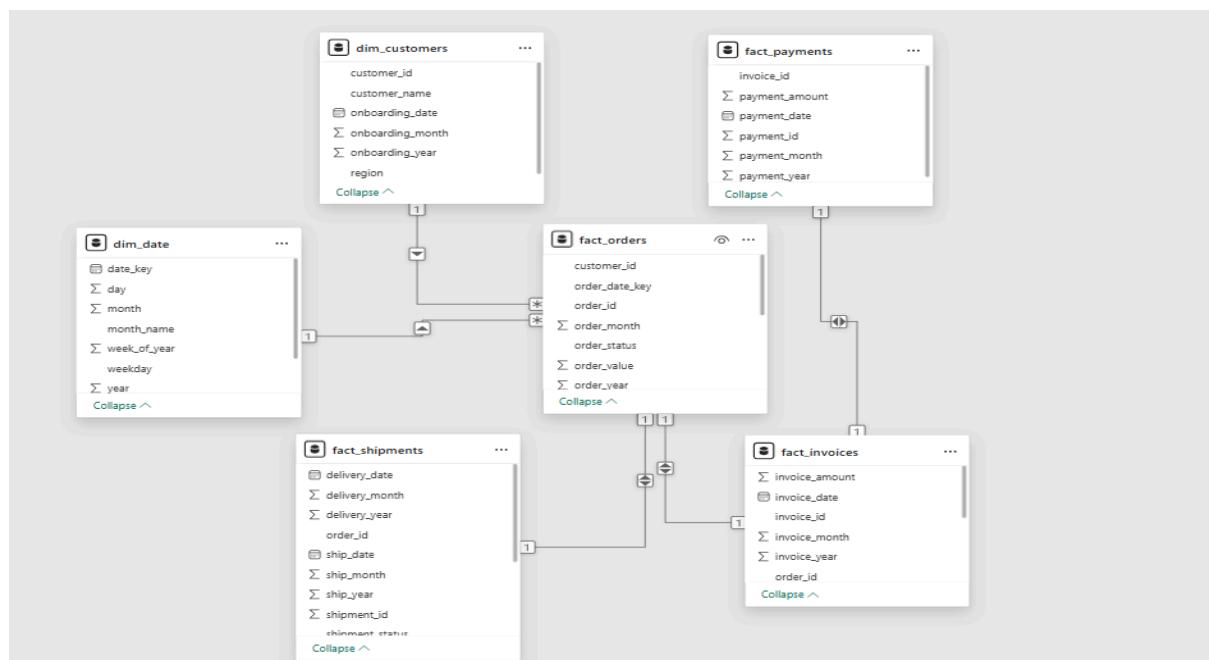
This warehouse layer enables analysis of revenue growth, order conversion, and payment collection efficiency.

Power BI Data Model

The warehouse is connected to Power BI and relationships are created to model the Order-to-Cash lifecycle.

Customer → Orders → Shipments → Invoices → Payments

The model allows tracking revenue from sale creation to final cash collection.



Business Metrics (KPIs)

The dashboard calculates business-level financial metrics:

- Total Revenue
- Collection Efficiency %
- Outstanding Amount
- Average Order Value (AOV)
- Repeat Customer Rate
- Customer Lifetime Value (CLV)
- Revenue at Risk
- Late Delivery Rate
- Aging Buckets (0-30, 31-60, 61-90, 90+ Days)

These metrics evaluate both **sales performance and financial health**.

A screenshot of a business intelligence dashboard interface. On the left, there is a sidebar with a dark header containing the text "Dashboard". Below the header, there is a list of items with small icons next to them, such as "Late Delivery %", "Order to Cash Days", "Orders YoY Growth", etc. One item, "Repeat Customer %", has a light gray rectangular highlight around it. To the right of the highlighted item is a small circular icon with a double-headed arrow and three dots. At the bottom of the sidebar, the text "Perspectives (0)" is displayed.

- Late Delivery %
- Order to Cash Days
- Orders YoY Growth
- Orders YTD
- Outstanding > 30 Days
- Outstanding Amount
- Repeat Customer %
- Repeat Customers
- Revenue at Risk
- Total Invoiced Amount
- Total Order Value
- Total Orders
- Total Paid Amount
- Total Shipments

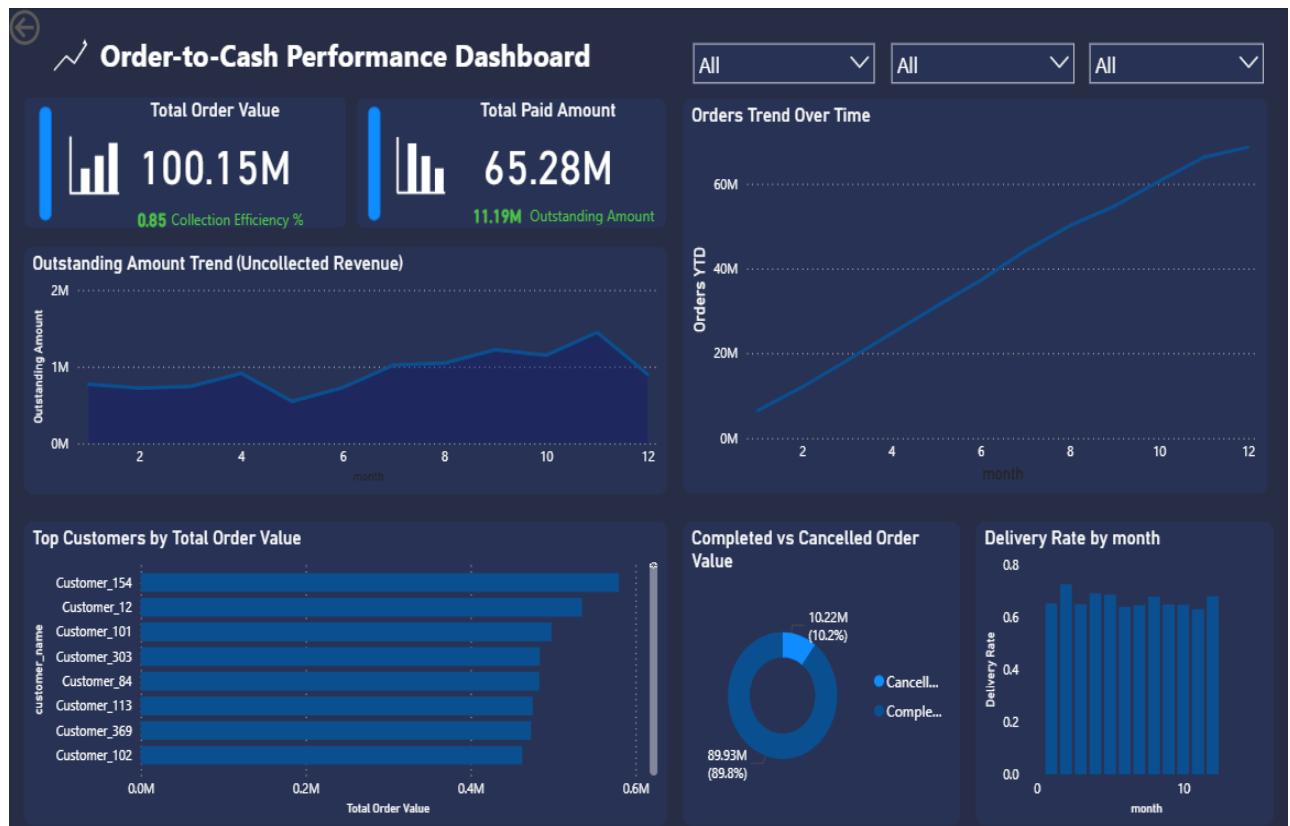
Perspectives (0)

Interactive Dashboard

An executive dashboard was built to monitor operational and financial performance.

It enables management to:

- Track revenue trends over time
- Monitor order growth
- Identify top customers
- Detect payment delays
- Evaluate collection efficiency
- Monitor overdue invoices



Key Insights Delivered

The dashboard helps stakeholders identify:

- Customers who generate high revenue but delay payments
- Revenue stuck in overdue invoices
- Delivery performance impact on payments
- Collection efficiency trends month-over-month

The system converts operational data into actionable financial intelligence.

Conclusion

This project demonstrates a complete data analytics workflow, not just visualization.

It covers:

Data ingestion → Data validation → Data warehousing → Data modeling → Business intelligence reporting

The platform provides a unified view of the company's revenue cycle and enables better financial decision-making by tracking how effectively revenue converts into cash.